

# Assignment-2 : Matrix Multiplication

MTCS-202(P) - Hadoop Assignment

Aryan Sai Arvapelly, Regd. No - 23352, I MTech CS

The **Hadoop Map-Reduce** program to calculate the product of two matrices,  $A(m \times n)$  and  $B(n \times p)$ , is as follows:

```
import java.io.IOException;
import org.apache.hadoop.conf.Configuration;
import org.apache.hadoop.fs.Path;
import org.apache.hadoop.io.IntWritable;
import org.apache.hadoop.io.Text;
import org.apache.hadoop.mapreduce.Job;
import org.apache.hadoop.mapreduce.Mapper;
import org.apache.hadoop.mapreduce.Reducer;
import org.apache.hadoop.mapreduce.lib.input.FileInputFormat;
import org.apache.hadoop.mapreduce.lib.output.FileOutputFormat;

public class MatrixMul {

    public static class MatrixMapper extends Mapper<Object,
Text, Text, Text> {
        public void map(Object key, Text value, Context context) throws IOException, InterruptedException {
            String[] tokens = value.toString().split(",");
            String matrix = tokens[0];
            int row = Integer.parseInt(tokens[1]);
            int col = Integer.parseInt(tokens[2]);
            int val = Integer.parseInt(tokens[3]);
```

```

        if (matrix.equals("A")) {
            for (int k = 0; k < context.getConfiguration().getInt("p", 1); k++) {
                context.write(new Text(row + "," + k),
                    new Text(matrix + "," + col + "," + val));
            }
        } else if (matrix.equals("B")) {
            for (int i = 0; i < context.getConfiguration().getInt("m", 1); i++) {
                context.write(new Text(i + "," + col),
                    new Text(matrix + "," + row + "," + val));
            }
        }
    }
}

public static class MatrixReducer extends Reducer<Text,
Text, Text, IntWritable> {
    public void reduce(Text key, Iterable<Text> values,
Context context) throws IOException, InterruptedException {
        int[] rowA = new int[context.getConfiguration().getInt("n", 1)];
        int[] colB = new int[context.getConfiguration().getInt("n", 1)];

        for (Text val : values) {
            String[] tokens = val.toString().split(",");

            String matrix = tokens[0];
            int index = Integer.parseInt(tokens[1]);
            int value = Integer.parseInt(tokens[2]);

            if (matrix.equals("A")) {
                rowA[index] = value;
            } else if (matrix.equals("B")) {
                colB[index] = value;
            }
        }
    }
}

```

```

        int result = 0;
        for (int i = 0; i < context.getConfiguration().
getInt("n", 1); i++) {
            result += rowA[i] * colB[i];
        }
        context.write(key, new IntWritable(result));
    }
}

public static void main(String[] args) throws Exception
{
    Configuration conf = new Configuration();
    conf.setInt("m", Integer.parseInt(args[2]));
    conf.setInt("n", Integer.parseInt(args[3]));
    conf.setInt("p", Integer.parseInt(args[4]));
    Job job = Job.getInstance(conf, "matrix multiplicat
ion");
    job.setJarByClass(MatrixMul.class);
    job.setMapperClass(MatrixMapper.class);
    job.setReducerClass(MatrixReducer.class);
    job.setOutputKeyClass(Text.class);
    job.setOutputValueClass(Text.class);
    FileInputFormat.addInputPath(job, new Path(args
[0]));
    FileOutputFormat.setOutputPath(job, new Path(args
[1]));
    System.exit(job.waitForCompletion(true) ? 0 : 1);
}
}


```

## Input

The program takes the input in the format: `matrix, row, column, value`

To run the program, we need to pass the input and output paths along with the dimensions of the two matrices: **m, n, p**.

```
hadoop jar MatrixMul.jar MatrixMul /matmul/input.txt /output m n p
```



## Mapper

The mapper receives input lines representing elements of Matrix A and Matrix B. For each input line, the mapper emits key-value pairs in the below way:

### 1. For Matrix A:

- **Intuition:** The value `val` at row `row` and column `col` in Matrix A needs to be multiplied with the value at row `row` and column `k` in Matrix B.
- The mapper receives an input value from Matrix A. It iterates over all columns of Matrix B for the given row of Matrix A.
- For each column `k`, it emits a key-value pair where the key is of the format `"row,k"` and the value is of the format `"A,col,val"`.

### 2. For Matrix B:

- **Intuition:** The value `val` at row `row` and column `col` in Matrix B needs to be multiplied with the value at row `i` and column `col` in Matrix A.
- The mapper receives an input value from Matrix B. It iterates over all rows of Matrix A for the given column of Matrix B.

- For each row `i`, it emits a key-value pair where the key is of the format `"i,col"` and the value is of the format `"B,row,val"`.

## Reducer

### 1. Array Initialization:

- Two integer arrays `rowA` and `colB` are initialized with the length equal to the number of columns in the matrices(`n`).

### 2. Adding elements to the arrays:

- For each received value, it parses the string representation into matrix identifier, index, and value.
- Based on the matrix identifier, the value is stored in the appropriate array (`rowA` for Matrix A and `colB` for Matrix B) at the corresponding index.

### 3. Matrix Multiplication:

- Perform the matrix multiplication by iterating over the arrays `rowA` and `colB`.
- For each index `i`, it computes the product of corresponding elements from `rowA` and `colB`, and accumulates the result.
- The accumulated result represents the value of the resulting matrix at the key (row, column) pair.

## Output

- The reducer emits the key-value pair where the key is the (row, column) pair and the value is the computed result of matrix multiplication at that position.

#### File contents

```
0,0 27
0,1 30
0,2 33
1,0 61
1,1 68
1,2 75
2,0 95
2,1 106
2,2 117
```